

# STAT 177, CLASS 5

Richard Waterman

July 2020

# OBJECTIVES

# OBJECTIVES

- Writing basic functions.
- Common data cleaning activities.
- More on the “groupby” command.
- From transactional to behavioral data sets.

# WRITING BASIC FUNCTIONS

# WRITING BASIC FUNCTIONS

- If you find yourself using almost identical blocks of code in different parts of your program, then you almost certainly want to write a “function”.
- A function helps you organize and reuse code.
- In math, a function  $f$  is a rule that takes an input  $x$  and provides a unique output,  $y$ .
- We write:  $y = f(x)$ .
- When we write Python code for a function, we have to define the inputs and create the rule that makes the output.
- Functions will be useful to us when we do data analysis, because we are often repeating the same activity on different data sets, or even the same activity within a dataset.
- These activities, could be data cleaning, merging, summarizing and so on.
- We will learn how to write a function, then later see how they can be applied to data.
- There is a special type of function in Python called a “lambda function” that is useful to create very simple functions.

# COMPONENTS OF A FUNCTION

- The Python keyword to create a function is 'def'.
- The keyword to end a function and return its result is 'return'.
- The body of a function needs to be indented.
- The function can take multiple inputs (parameters/arguments).
- These inputs can be specified as either positional inputs or keyword inputs.
- The keyword inputs are typically used for those arguments that take default values or are optional.
- We will write some functions to illustrate the structure.
- The function is *called* (run) by using its name and then parentheses().

# DEGREES CELSIUS TO FAHRENHEIT

- We will write a function, called 'tconvert', that takes a single input and converts Celsius to Fahrenheit.
- The conversion formula is  $f(c) = 32 + (9/5)c$ , where  $c$  is the input temperature.

```
def tconvert(x):  
    y = 32 + 9/5 * x  
    return y  
  
# Call the function a couple of times  
print(tconvert(22))  
print(tconvert(100))
```

```
71.6  
212.0
```

# ANATOMY OF THE FUNCTION

The diagram illustrates the components of a Python function definition. It shows the code for a function named `tconvert` that takes a parameter `x` and returns a value `y`. The code is annotated with numbers 1 through 7 and arrows pointing to specific parts:

```
1  def
2  tconvert
3  (
4  x):
5  y = 32 + 9/5 * x
6  return
7  y
```

Annotations:

- 1: Points to the `def` keyword.
- 2: Points to the function name `tconvert`.
- 3: Points to the opening parenthesis `(`.
- 4: Points to the parameter `x`.
- 5: Points to the assignment statement `y = 32 + 9/5 * x`.
- 6: Points to the `return` keyword.
- 7: Points to the variable `y` being returned.



## ANATOMY OF THE FUNCTION (CTD.)

1. The keyword “def” signals a function is about to be defined.
2. This particular function is named *tconvert* .
3. The *parameter* to be input is called “x”.
4. The colon indicates we are about to get an indented code block.
5. The *body* of the function.
6. The keyword “return” statement.
7. The value to be returned by the function.

# ADDING A DOCSTRING TO THE FUNCTION

- The docstring describes the action of the function.
- If it is there, then it is the first string right after the function definition.
- Using three quotes, `'''`, allows you the potential to create a multiline string.
- You can print the docstring using the `'__doc__'` method as below.

```
def tconvert(x):  
    '''Takes in a number x, returns the converted temperature of x.'''  
    y = 32 + 9/5 * x  
    return y  
  
print(tconvert.__doc__)
```

```
Takes in a number x, returns the converted temperature of x.
```

# CHECKING THE INPUT TYPE

- Look what happens if we call this function on a string input:

```
print(tconvert("It's hot!"))
```

-----  
TypeError

Traceback (most recent call last)

<ipython-input-3-bbb8240b6a68> in <module>

----> 1 print(tconvert("It's hot!"))

<ipython-input-2-fcc95a68b6ad> in tconvert(x)

1 def tconvert(x):

2 '''Takes in a number x, returns the converted temperature of x.'''

----> 3 y = 32 + 9/5 \* x

4 return y

5

# CHECKING THE INPUT TYPE

- It makes sense to only have the function work if it has a numeric input.
- The input type can be checked by using the “isinstance” command.

```
def tconvert(x):  
    '''Takes in a number x, returns the converted temperature of x.'''  
    if( isinstance(x, (int, float)) and not isinstance(x, (bool))): #Booleans are also ints!  
        y = 32 + 9/5 * x  
        return y  
    else:  
        return "This function only works with numbers."  
  
print(tconvert(3))  
print(tconvert(21.2))  
print(tconvert(True))  
print(tconvert("It's hot!"))
```

```
37.4  
70.16  
This function only works with numbers.  
This function only works with numbers.
```

# MULTIPLE PARAMETERS

- We will give the function an additional argument, that allows the possibility of going from 'f2c' or 'c2f'.
- We check that it takes on one of the two valid values.

```
def tconvert(x, direction):  
    '''Takes in a number x, returns the converted temperature of x.'''  
  
    if not ( (direction == "f2c") or (direction == "c2f")):  
        return "Direction must be 'c2f' or 'f2c'."  
    if not ( isinstance(x, (int, float)) and not isinstance(x, (bool))) : # Booleans are also ints!  
        return "This function only works with numbers."  
  
    if direction == "c2f":  
        y = 32 + 9/5 * x  
        return y  
    if direction == "f2c":  
        y = (x - 32) * 5/9  
        return y
```

# USING THE IMPROVED FUNCTION

```
print(tconvert(0.0, "c2f"))  
print(tconvert(71, "f2c"))  
print(tconvert(71, "a2b")) # Break it on purpose.  
print(tconvert(225, "c2f"))
```

```
32.0  
21.666666666666668  
Direction must be 'c2f' or 'f2c'.  
437.0
```

# KEYWORD ARGUMENTS

- You can also provide arguments with the key = value syntax.
- This way, if there are multiple arguments, the order of the arguments does not matter.

```
print(tconvert(direction = "f2c", x = 21))
```

```
-6.111111111111111
```

# KEYWORD ARGUMENTS CAN'T COME BEFORE POSITIONAL ARGUMENTS

- There are rules for the order of the arguments.

```
print(tconvert(direction = "f2c", 212)) # This is illegal.
```

```
File "<ipython-input-8-5f11852d7b81>", line 1
  print(tconvert(direction = "f2c", 212)) # This is illegal.
                        ^
```

```
SyntaxError: positional argument follows keyword argument
```



# DEFAULT VALUES

- When a function has many arguments it makes sense to give the less used ones, default values.
- Using the 'parameter = value' syntax will do this.
- We will assume that the user wants to go from c2f, unless they say otherwise.
- Note the 'direction = "c2f"' in the argument list below.

```
def tconvert(x, direction = "c2f"):
    '''Takes in a number x, returns the converted temperature of x.'''

    if not ( (direction == "f2c") or (direction == "c2f") ):
        return "Direction must be 'c2f' or 'f2c'."
    if not ( isinstance(x, (int, float)) and not isinstance(x, (bool)) ): # Booleans are also ints!
        return "This function only works with numbers."

    if direction == "c2f":
        y = 32 + 9/5 * x
        return y
    if direction == "f2c":
        y = (x - 32) * 5/9
        return y
```

# CALLING THE FUNCTION WITH THE DEFAULT PARAMETER

```
print(tconvert(10)) # No need for the direction parameter, if happy with 'c2f'  
print(tconvert(200))
```

```
50.0  
392.0
```

## RETURNING MORE THAN ONE VALUE

- You can return multiple values from a function as a ‘tuple’.
- The function below takes a pandas series and returns the top 3 occurring levels.
- We make some data by randomly sampling letters of the alphabet.

```
def top3(x):  
    y = x.value_counts(sort=True) # Recall that value_counts makes frequencies, and the sort argument  
    will sort them.  
    return y.index[0], y[0], y.index[1], y[1], y.index[2], y[2] # Return multiple values.
```

```
import pandas as pd  
import numpy as np  
np.random.seed(1234)  
data = np.random.choice(['a','b','c','d','e','f','g','h','i','j'], size=1000, replace=True) # Random  
    sampling with replacement.  
data = pd.Series(data) # Store the data in a pandas Series.
```

# RETURNING MULTIPLE VALUES

- Note how the results come back as a tuple.

```
top3(data)
```

```
('j', 116, 'd', 108, 'b', 107)
```

# LAMBDA FUNCTIONS

- These are special functions, with no name (anonymous), that are defined in an expression rather than a statement.
- They are suitable for very simple functions that are used for a short period of time.
- They are useful in data analysis, because we are often trying different transforms of the data, and we can make a function, which itself takes another function (the lambda function) as an argument.

```
import math
def data_transform(x, fn):
    return [fn(input) for input in x] # Notice that the function "fn" is being used in the
    comprehension.

data = [1,2,5,6,2,1]

print(data_transform(data, lambda y: y**2))
print(data_transform(data, lambda y: math.log(y)))
```

```
[1, 4, 25, 36, 4, 1]
[0.0, 0.6931471805599453, 1.6094379124341003, 1.791759469228055, 0.6931471805599453, 0.0]
```

# THE ‘MAP’ COMMAND

- Another place you might see lambda functions used is with the “map” command.
- ‘map’ will map a function to the elements of a container.
- You may see ‘map’ in someone else’s Python code, but many prefer a list comprehension instead.

```
data = np.random.randint(10, size=10000) # Make some data, random integers between 0 and 9.
```

```
list(map(lambda x: x**2 , data))[0:5] # Just look at the first 5 elements.
```

```
[25, 49, 0, 4, 9]
```

## MORE ABOUT FUNCTIONS

- There's more to learn about functions, such as 'Arbitrary Arguments' (\*args) and 'Arbitrary Keyword Arguments' (\*\*kwargs), but we will discuss these as they become necessary.

# COMMON DATA CLEANING ACTIVITIES



# COMMON DATA CLEANING ACTIVITIES

- Truth be told, what is *common* most likely depends on the type of data you are used to.
- Cleaning a time-series could be quite different from cleaning a list of words.
- Examples of cleaning/pre-processing activities include:
  - There are rows in the data frame that need to be removed.
  - There is missing data that needs to be addressed.
  - There are typos in the data that need to be corrected.
  - Things that should be numeric are held as strings.
  - Dates have not been properly parsed.
  - Levels of a categorical variable need to be collapsed.
  - Sorting data.

# REMOVING ROWS FROM THE DATA FRAME

- We will read in a data frame and remove rows.
  - Remove the first row.
  - Remove the last row.
  - Remove rows using a logical filter.

```
# Read in a data frame for illustration.
import os
os.chdir('C:\\Users\\richardw\\Dropbox (Penn)\\Teaching\\477s2020\\DataSets')
op_data = pd.read_csv("Outpatient_to_clean.csv")
print(op_data.shape) # Track the dimensions.
op_data.drop(0, inplace=True) # Drop the first row.
print(op_data.shape) # Track the dimensions.
op_data.drop(len(op_data)-1, inplace=True) # Drop the last row.
print(op_data.shape) # Track the dimensions.
```

```
(3699, 9)
(3698, 9)
(3697, 9)
```

# REMOVE ROWS USING A LOGICAL FILTER

```
temp = op_data.loc[op_data['Status'] != "Bumped"] # The logical filter selects all rows, that are not Bumped.  
temp['Status'].value_counts() # Note that "Bumped has gone".
```

```
Arrived      2163  
Cancelled    796  
No Show      526  
Rescheduled   1  
Name: Status, dtype: int64
```

## USING THE READ\_CSV ARGUMENTS

- If you know upfront that that certain rows at the beginning or end of the data frame need to be dropped, you could use the arguments to `read_csv`:
- You can specify the specific line numbers to skip, or number of lines at the bottom of the file to skip with the arguments:
  - `skiprows`
  - `skipfooter`



1		HISPANIC	Cancelled
2	AFRICAN	AMERICAN	NaN
3		HISPANIC	Bumped
4		HISPANIC	No Show

# USING .DROPNA()

- The .dropna() would be a common pre-processing step.

```
op_data.dropna(inplace=True) # '.dropna()' removes all rows with any missing data.
print(op_data.head(5)) # Note that some of the rows have disappeared.
```

	PID	SchedDate	ApptDate	Dept	Language	Sex	Age	\
0	P10092	2012-07-27	2012-10-05	DERM	ENGLISH	F	80+	
1	P10151	2013-11-28	2014-01-03	PULMONARY	SPANISH	Didn't say	32	
6	P10410	2013-10-31	2013-11-03	ORTHOPAEDICS	ENGLISH	M	54	
10	P10391	2012-12-24	2012-12-27	GENERAL SURGERY	ENGLISH	F	49	
12	P10138	2011-11-14	2011-11-20	CPO	ENGLISH	F	45	
		Race	Status					
0		AFRICAN AMERICAN	Arrived					
1		HISPANIC	Cancelled					
6		AFRICAN AMERICAN	No Show					
10		WHITE (NON-HISPANIC)	Arrived					
12		HISPANIC	Arrived					

# TYPOS AND STRINGS

- If you have a typo, then there is the possibility to make it missing, or to overwrite the bad values if you know what they should be.
- In the outpatient data all values for Sex should be either ‘M’ or ‘F’
- We will look through the Sex column, changing elements that are neither ‘M’ or ‘F’ to ‘Unknown’.
- If the .map method doesn’t find a key in the dict structure, for the element, then it replaces it with NA.
- The .fillna, then replaces these NAs with ‘Unknown’.

```
op_data['Sex'] = op_data['Sex'].map({'F': 'F', 'M': 'M'}).fillna('Unknown')
print(op_data[0:7])
```

	PID	SchedDate	ApptDate	Dept	Language	Sex	Age	\
0	P10092	2012-07-27	2012-10-05	DERM	ENGLISH	F	80+	
1	P10151	2013-11-28	2014-01-03	PULMONARY	SPANISH	Unknown	32	
6	P10410	2013-10-31	2013-11-03	ORTHOPAEDICS	ENGLISH	M	54	
10	P10391	2012-12-24	2012-12-27	GENERAL SURGERY	ENGLISH	F	49	
12	P10138	2011-11-14	2011-11-20	CPO	ENGLISH	F	45	
13	P10677	2011-04-17	2011-05-09	GASTROENTEROLOGY	ENGLISH	Unknown	31	
14	P10229	2013-10-07	2013-11-18	VASCULAR SURGERY	ENGLISH	F	39	



	Race	Status
0	AFRICAN AMERICAN	Arrived
1	HISPANIC	Cancelled
6	AFRICAN AMERICAN	No Show
10	WHITE (NON-HISPANIC)	Arrived
12	HISPANIC	Arrived

# REMOVING WHITE SPACE

- In this data frame the elements of the ‘Dept’ column are inconsistent.
- Some of them start/end with blank spaces.
- We will remove all white space from the strings in ‘Dept’ using a method called `.str.strip()`.

```
op_data['Dept'] = op_data['Dept'].str.strip()
print(op_data['Dept'])
```

```
0          DERM
1    PULMONARY
6    ORTHOPAEDICS
10   GENERAL SURGERY
12          CPO
...
3694   RHEUMATOLOGY
3695   OTOLARYNGOLOGY
3696   VASCULAR SURGERY
3697   NEUROLOGICAL
3698          DERM
Name: Dept, Length: 3688, dtype: object
```

# DATA THAT SHOULD BE NUMERIC ARE HELD AS STRINGS

- In this dataset there is an Age value recorded as '80+'.
- This is enough for the columns to be read in as a string.
- The `to_numeric` function will convert the strings to numbers if possible, but if not, replace them with NA.

```
print(op_data[0:2])
op_data['Age'].mean() # This fails.
```

	PID	SchedDate	ApptDate	Dept	Language	Sex	Age	\
0	P10092	2012-07-27	2012-10-05	DERM	ENGLISH	F	80+	
1	P10151	2013-11-28	2014-01-03	PULMONARY	SPANISH	Unknown	32	

  

		Race	Status
0	AFRICAN	AMERICAN	Arrived
1		HISPANIC	Cancelled

-----

ValueError

Traceback (most recent call last)

# CONVERTING FROM STRING TO NUMERIC

- The “errors = ‘coerce’” argument below will replace the non-coercible elements with NAs.

```
op_data['Age'] = pd.to_numeric(op_data['Age'], errors = 'coerce')  
print(op_data['Age'].mean()) # Now we can calculate the mean.
```

```
45.64065592309867
```

# DATES HAVE NOT BEEN PROPERLY PARSED

- As we saw in the class 4 notes, there are formatting methods, that you can use to control exactly how to parse a string as a date.
- See Class 4 notes, 'pd.to\_datetime', slides 7.5 and 7.7.

# LEVELS OF A CATEGORICAL VARIABLE NEED TO BE COLLAPSED

- This is a very common task and there will be many ways to do it.
- One that we have already seen is to use the `.map` method.
- But if there are many levels this could be cumbersome.
- In this data set there are two 'Dept' levels, NEUROLOGICAL and NEUROLOGY, that should be the same level.
- We can search for one, and replace it with the other:

# THE ORIGINAL FREQUENCY DISTRIBUTION

```
print(op_data['Dept'].value_counts())
```

NEUROLOGICAL	668
CPO	657
ORTHOPAEDICS	381
DERM	236
PLASTIC SURGERY	221
GENERAL SURGERY	179
UROLOGY	177
GASTROENTEROLOGY	162
PODIATRY	161
VASCULAR SURGERY	124
NEPHROLOGY	122
RHEUMATOLOGY	116
OTOLARYNGOLOGY	114
PULMONARY	108
TRAUMA	89

# REPLACING A SINGLE LEVEL

- We use the ‘.loc’ method, creating a logical filter for the rows we want to change.

```
op_data.loc[op_data['Dept'] == "NEUROLOGICAL", 'Dept'] = "NEUROLOGY" # The logical replacement filter.  
op_data['Dept'].value_counts() # Now there are more rows in the NEUROLOGY column.
```

NEUROLOGY	711
CPO	657
ORTHOPAEDICS	381
DERM	236
PLASTIC SURGERY	221
GENERAL SURGERY	179
UROLOGY	177
GASTROENTEROLOGY	162
PODIATRY	161
VASCULAR SURGERY	124
NEPHROLOGY	122
RHEUMATOLOGY	116
OTOLARYNGOLOGY	114
PULMONARY	108
TRAUMA	89



# SORTING A DATA FRAME

- You can sort by the Index (rownames) or by the values in a column.
- You can also sort on multiple columns.

```
op_data.sort_values(by='Age', inplace = True)
print(op_data.head(3))
```

	PID	SchedDate	ApptDate		Dept	Language	Sex	Age	\
3294	P10453	2013-05-22	2013-05-22	PEDIATRIC	SURGERY	ENGLISH	M	0.0	
821	P10453	2013-05-20	2013-05-22	PEDIATRIC	SURGERY	ENGLISH	M	0.0	
2546	P10993	2014-01-14	2014-01-20	PEDIATRIC	SURGERY	ENGLISH	M	0.0	
		Race	Status						
3294	AFRICAN	AMERICAN	Arrived						
821	AFRICAN	AMERICAN	Cancelled						
2546	AFRICAN	AMERICAN	Arrived						

# SORT BY PATIENT PID, THEN SCHEDULE DATE WITHIN PID

```
op_data.sort_values(by = ['PID', 'SchedDate'], inplace=True) # Sorting by two columns.
print(op_data.head(5))
```

	PID	SchedDate	ApptDate	Dept	Language	Sex	Age	\
2598	P10001	2012-11-23	2012-11-26	NEUROLOGY	ENGLISH	M	45.0	
499	P10002	2011-12-02	2011-12-21	NEUROLOGY	ENGLISH	M	11.0	
2733	P10002	2011-12-21	2012-06-20	NEUROLOGY	ENGLISH	M	11.0	
1618	P10003	2013-10-18	2014-02-16	PLASTIC SURGERY	ENGLISH	M	11.0	
3043	P10004	2012-09-26	2012-10-08	NEUROLOGY	BENGALI	M	39.0	
		Race		Status				
2598	WHITE	(NON-HISPANIC)		Arrived				
499	WHITE	(NON-HISPANIC)		Arrived				
2733	WHITE	(NON-HISPANIC)		Cancelled				
1618		AFRICAN AMERICAN		Arrived				
3043		OTHER		Arrived				

# THE GROUPBY COMMAND

# THE GROUPBY COMMAND

- Summarizing a data set by the levels of a categorical variable is an essential activity.
- We saw 'groupby' used with `.value_counts()` to summarize over the day-of-week or month-of-year.
- But it is very flexible and can be used to summarize in a more sophisticated way.
- In our outpatient dataset we may wish to summarize by patient ID (PID) to create a behavioral view of the patient.
- We will start by counting the number of appointments each patient has had.

# TRANSACTIONS TO BEHAVIOR

- Many databases are designed to capture *transactions* .
- But many analytic techniques are applied to behavioral patterns.
- This requires individual transactions to be aggregated to behaviors via a common ID.

# EXAMPLES

- Hospital:
  - The patient admissions database shows you who is in the hospital.
  - Combining admissions gives you a patient's history.
- Supermarket:
  - Individual items get aggregated to a customer's shopping visit.
  - Customer visits get aggregated to a long term behavior.
- HR:
  - A cut of the employee database in any given month, shows you who is employed.
  - Combining the months, shows you an employee's history.
- Outpatient clinic:
  - The visits database tells you about each visit.
  - Combining the visits, shows you a patient's history.

# THE COMMONALITY

- All of these examples are of the same essential nature.
- Transactions are recorded.
- A common ID can be used to link transactions.
- Histories can be built across common IDs.
- This translates transactions to {} behavior/history.

# COUNTING PATIENT VISITS

- groupby and using the count method will tell us how many appointments each patient had.

```
op_data.groupby('PID')['PID'].count()
```

```
PID
P10001      1
P10002      2
P10003      1
P10004      4
P10005      1
      ..
P10996      4
P10997     10
P10998     29
P10999      2
P11000      5
Name: PID, Length: 996, dtype: int64
```



# PRIOR VISITS

- If we are to use the behavioral history as a predictor, we need to use prior visits, not including the latest one.

```
op_data.groupby('PID')['PID'].count() - 1 # Remove 1 from the count.
```

```
PID
P10001    0
P10002    1
P10003    0
P10004    3
P10005    0
..
P10996    3
P10997    9
P10998   28
P10999    1
P11000    4
Name: PID, Length: 996, dtype: int64
```

# GROUPBY AND AGGREGATE

- We can summarize more than one variable at a time.
- The .agg method lets you create a new set of columns, based on a function applied to the groups.
- You can define this function yourself.
- Below we count the number of prior visits, and the number of prior visits that were of type “Arrived”.

```
prev = lambda x: len (x) - 1 # A lambda function to count previous appointments.
prior_arrived = lambda x: sum(x[:-1] == "Arrived") if len(x) > 1 else 'NA' # Counting the number of
prior arrivals.

op_behave = op_data.groupby('PID').agg(PriorVisits=('PID', prev),
                                      PriorArrived=('Status', prior_arrived))
print(op_behave.head(3))
```

	PriorVisits	PriorArrived
PID		
P10001	0	NA
P10002	1	1
P10003	0	NA

# KEEPING THE FEATURES OF THE MOST RECENT VISIT

- We still want to make sure we have the features of the current visit too.
- The ‘.last’ method will pull off the ‘last’ (most recent) visit.

```
op_last_values = op_data.groupby('PID').last() # Obtain the last row for each patient
print(op_last_values.head(5))
```

PID	SchedDate	ApptDate	Dept	Language	Sex	Age	\
P10001	2012-11-23	2012-11-26	NEUROLOGY	ENGLISH	M	45.0	
P10002	2011-12-21	2012-06-20	NEUROLOGY	ENGLISH	M	11.0	
P10003	2013-10-18	2014-02-16	PLASTIC SURGERY	ENGLISH	M	11.0	
P10004	2013-05-19	2013-05-26	PODIATRY	BENGALI	M	40.0	
P10005	2013-06-07	2013-07-29	DERM	ENGLISH	F	24.0	

  

PID	Race	Status
P10001	WHITE (NON-HISPANIC)	Arrived
P10002	WHITE (NON-HISPANIC)	Cancelled
P10003	AFRICAN AMERICAN	Arrived
P10004	OTHER	Arrived
P10005	WHITE (NON-HISPANIC)	No Show

# CHECKING THE CODE

- It is always a good idea to check your work.
- Here we pull off one patient and can then compare to the original sorted .csv file, to make sure it seems right.

```
op_last_values.loc['P10998']
```

```
SchedDate    2013-02-22 00:00:00
ApptDate     2013-03-02 00:00:00
Dept         CPO
Language     ENGLISH
Sex          M
Age         60
Race        AFRICAN AMERICAN
Status       Arrived
Name: P10998, dtype: object
```

# MERGE THE TWO DATA FRAMES

- We finish by merging the two new data frames on the PID, to create a single data set for analysis.

```
final_op_data = pd.merge(op_last_values, op_behave, on = 'PID')
print(final_op_data.head(5))
```

	SchedDate	ApptDate	Dept	Language	Sex	Age	\
PID							
P10001	2012-11-23	2012-11-26	NEUROLOGY	ENGLISH	M	45.0	
P10002	2011-12-21	2012-06-20	NEUROLOGY	ENGLISH	M	11.0	
P10003	2013-10-18	2014-02-16	PLASTIC SURGERY	ENGLISH	M	11.0	
P10004	2013-05-19	2013-05-26	PODIATRY	BENGALI	M	40.0	
P10005	2013-06-07	2013-07-29	DERM	ENGLISH	F	24.0	
		Race	Status	PriorVisits	PriorArrived		
PID							
P10001	WHITE (NON-HISPANIC)	Arrived		0	NA		
P10002	WHITE (NON-HISPANIC)	Cancelled		1	1		
P10003	AFRICAN AMERICAN	Arrived		0	NA		
P10004	OTHER	Arrived		3	2		
P10005	WHITE (NON-HISPANIC)	No Show		0	NA		

# SUMMARY

# SUMMARY

- Writing basic functions.
- Common data cleaning activities.
- More on the “groupby” command.
- From transactional to behavioral data sets.

# NEXT TIME



# NEXT TIME

- Graphics